



# Managing AWS Infrastructure with Terraform v1

# Agenda




- Iterative / incremental Infrastructure-as-Code
- Basics
- Templating
- Workspaces
- Samples
- Optional: hands-on
- Q & A


# Iterative / incremental Infrastructure-as-Code

t=0: template version 1 created

Terraform Template



AWS Resources




Terraform Workspace (Object in S3)




t=1 : template version 1 applied (physical resources created)

Terraform Template



AWS Resources

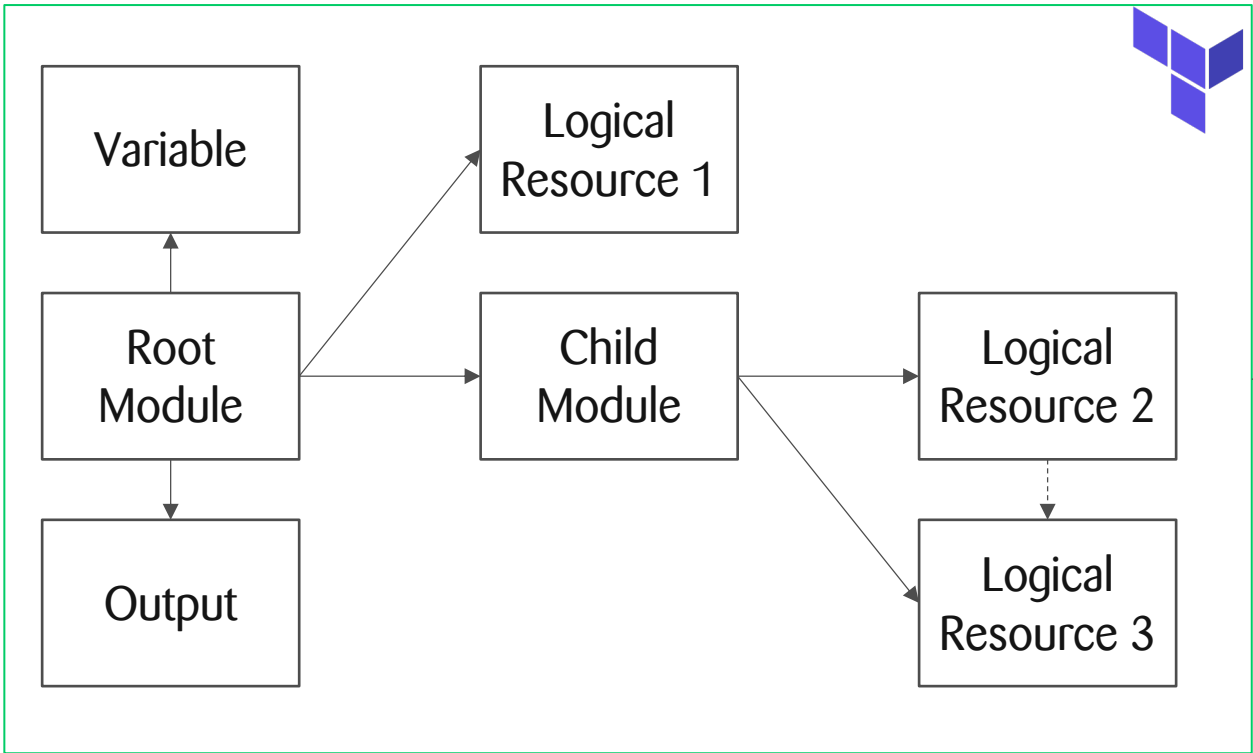


Terraform Workspace (Object in S3)




t=2: template version 2 created (logical resources added)

Terraform Template



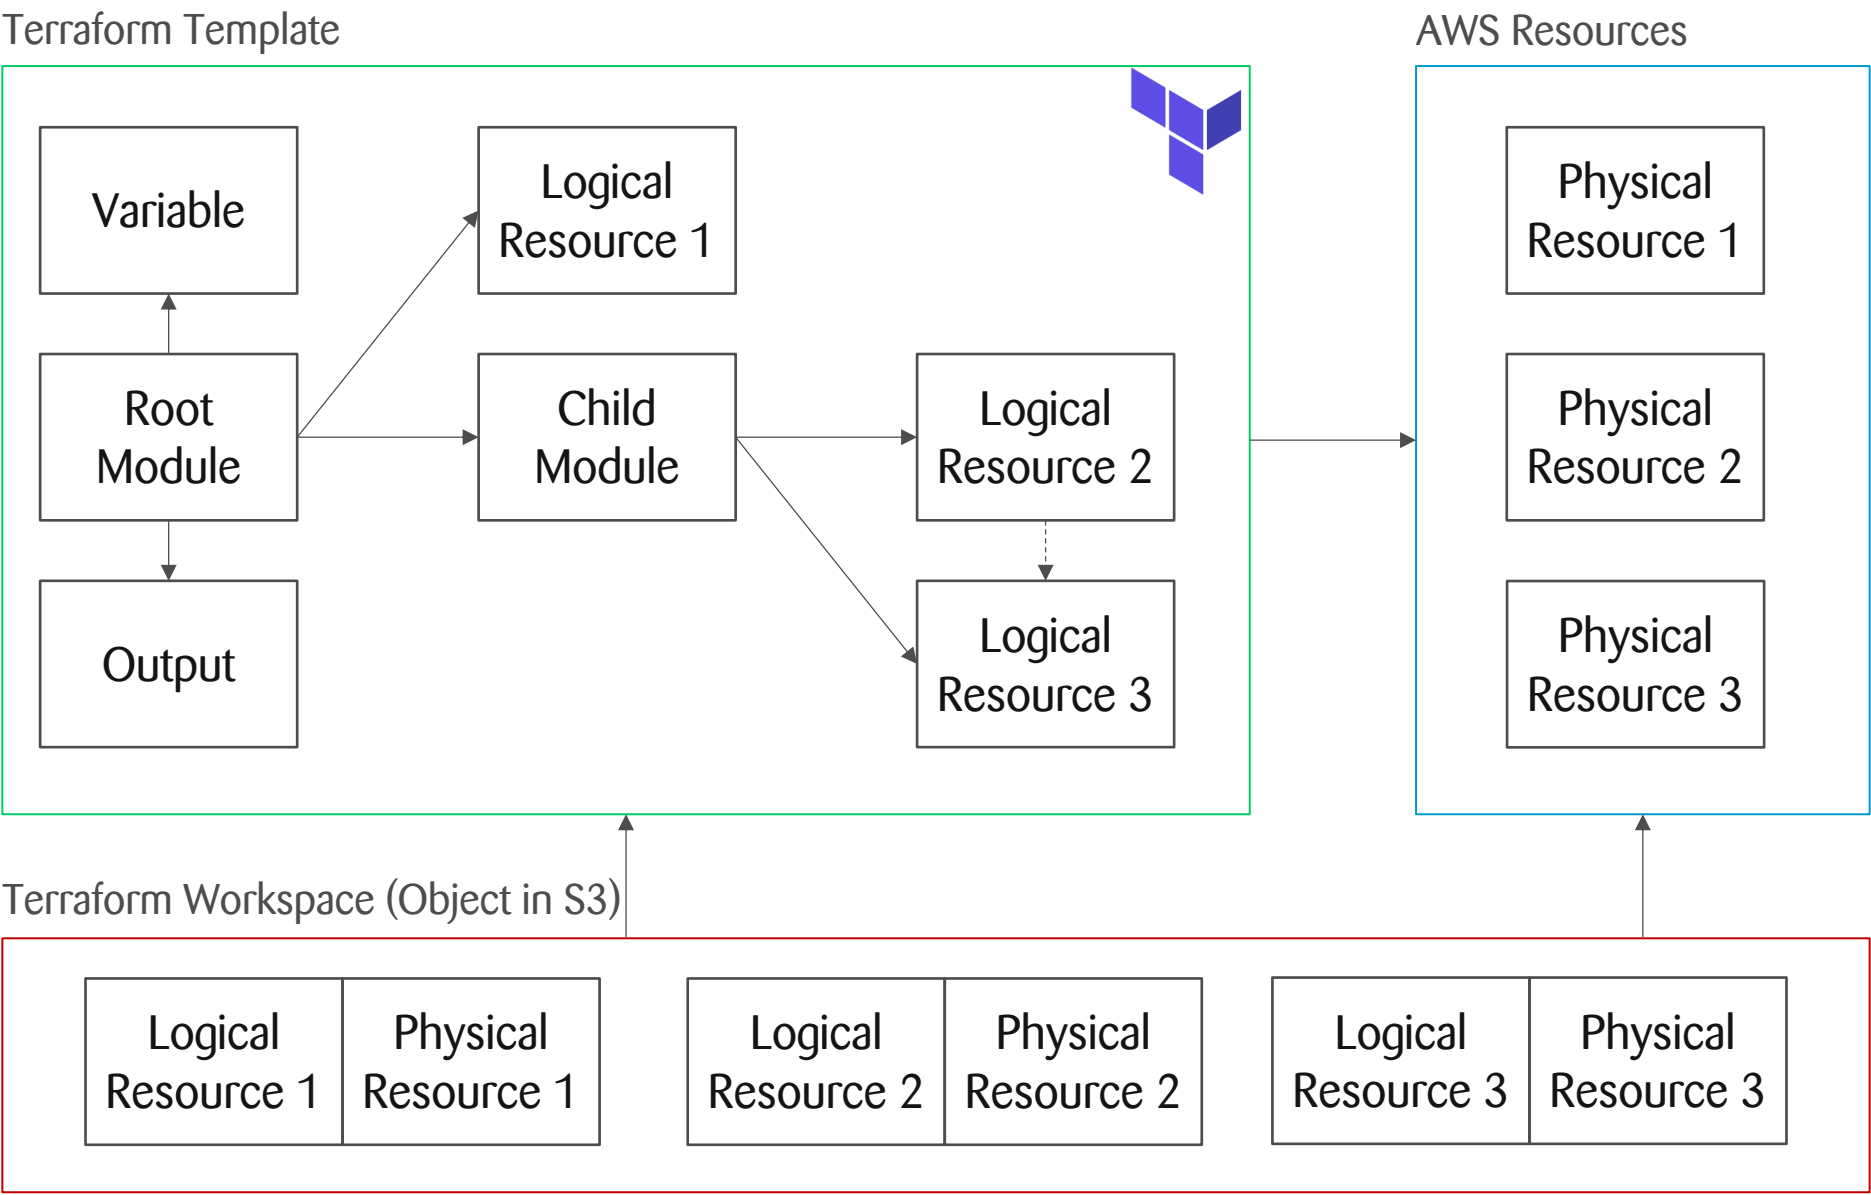
AWS Resources



Terraform Workspace (Object in S3)

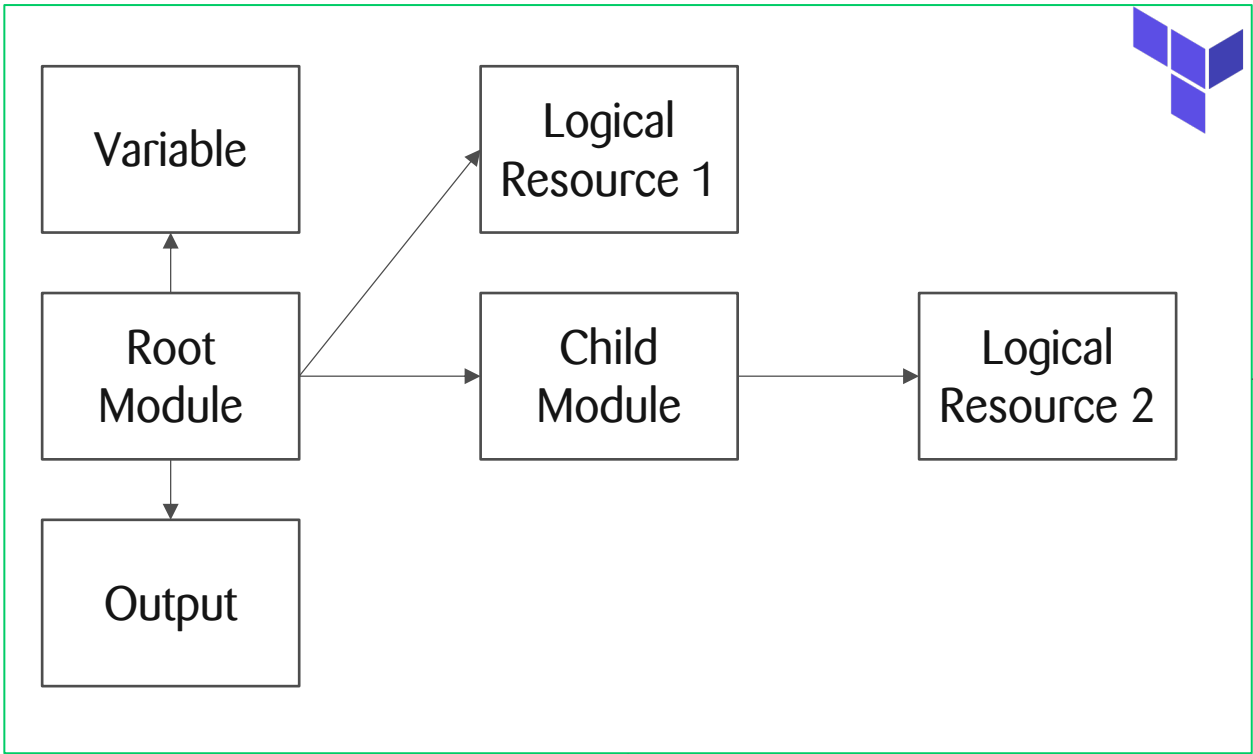


t=3: template version 2 applied (physical resources added)




t=4: template version 3 created (logical resource removed)

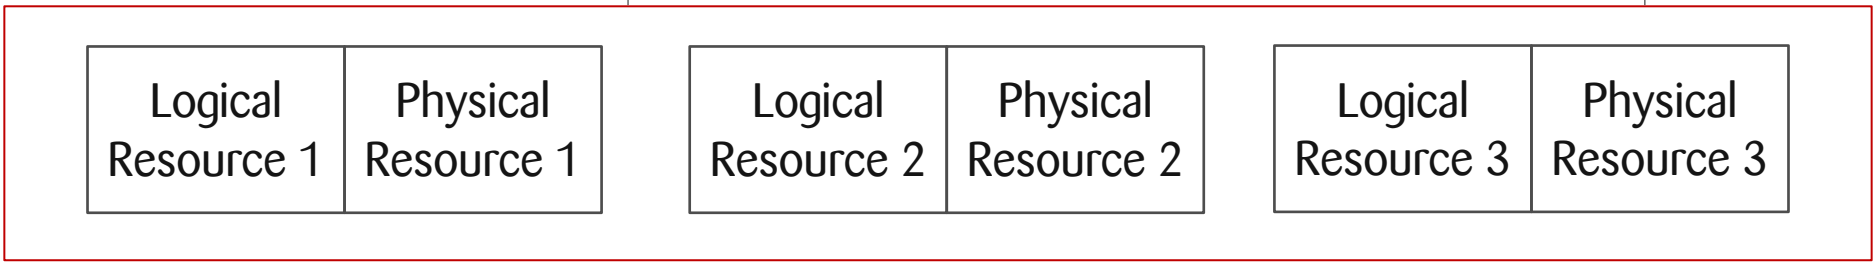
Terraform Template



AWS Resources




Terraform Workspace (Object in S3)



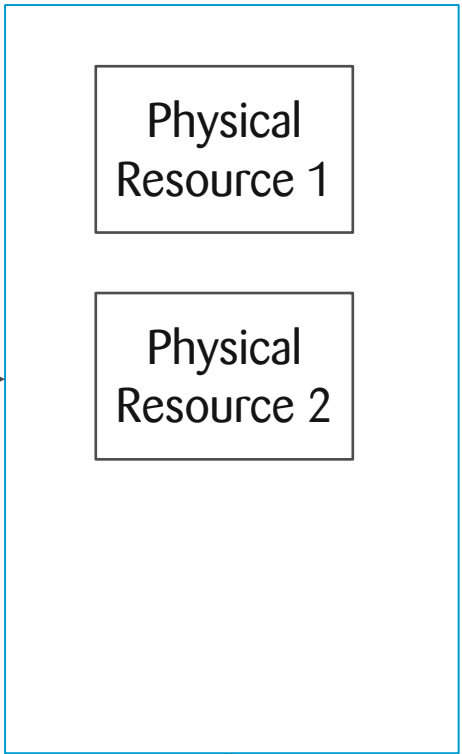


t=5: template version 3 applied (physical resource removed)

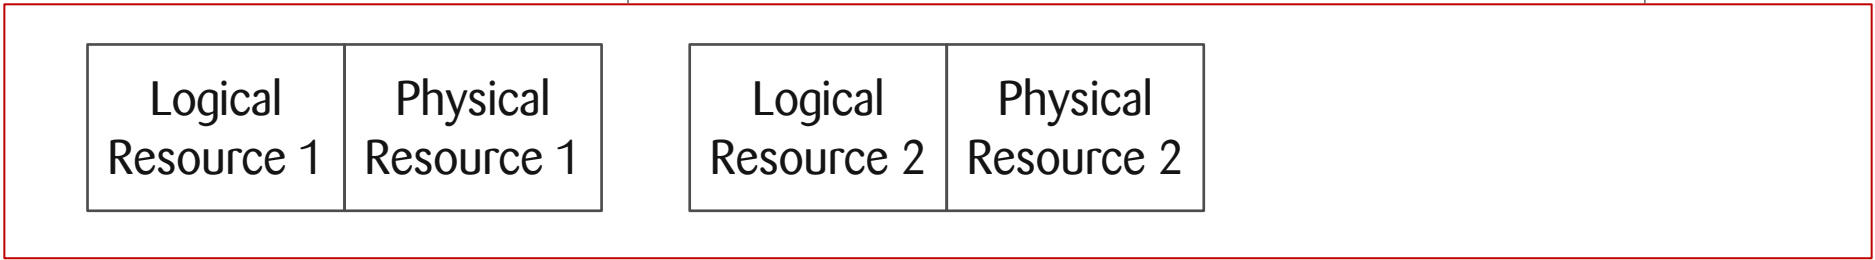
Terraform Template



AWS Resources



Terraform Workspace (Object in S3)



# Basics

# Terraform

## Module Layout (Filesystem)



### Sample (recommended structure and filenames for a minimal module)

```
Directory: RootModuleA
```

```
File: variables.tf
```

```
File: main.tf
```

```
File: outputs.tf
```

```
Directory: ./modules/ChildModule1
```

```
File: variables.tf
```

```
File: main.tf
```

```
File: outputs.tf
```

```
Directory: ./modules/ChildModule<n>
```

```
...
```

```
Directory: RootModule<n>
```

```
...
```

# Terraform

## Provider



- The Amazon Web Services (AWS) provider is used to interact with the many resources supported by AWS
- The provider needs to be configured with the proper credentials before it can be used (static credentials, environment variables, shared credentials file, EC2 Role)

```
provider "aws" {  
    version = "~> 2.0"  
    region  = "eu-central-1"  
}
```

# Terraform

## Backend



- Terraform must store state about your managed infrastructure and configuration. This state is used by Terraform to map real world resources to your configuration
- A "backend" in Terraform determines how state is loaded
- By default, Terraform uses the "local" backend
- The S3 backend stores the state as a given key in a given bucket on Amazon S3

```
terraform {  
  backend "s3" {  
    bucket = "terraformchtzbucket"  
    key     = "bucketsample/terraform.tfstate"  
    region = "eu-central-1"  
  }  
}
```

# Terraform

## Module



- A module is a container for multiple resources that are used together
- A module consists of the resources defined in the .tf files in the module directory
- Every Terraform configuration has at least one module, known as its root module
- A module can call other modules, which lets you include the child module's resources into the configuration.
- Modules can also be called multiple times, either within the same configuration or in separate configurations

```
module "bucket_b" {  
    source = "../mybucket"  
    bucket_suffix = "chtz-testbucket-b"  
}
```

- Each resource block describes one or more infrastructure objects, such as virtual networks, compute instances, or higher-level components such as DNS records

```
resource "aws_s3_bucket" "samplebucket" {  
    bucket = "${terraform.workspace}-${var.bucket_suffix}"  
}
```

```
resource "aws_s3_bucket_public_access_block" "samplebucket_nonpublic" {  
    count = var.block_public_acls ? 1 : 0  
    bucket = aws_s3_bucket.samplebucket.id  
    block_public_acls = true  
}
```

# Terraform

## Data Sources



- Data sources allow data to be fetched or computed for use elsewhere in Terraform configuration
- Each provider may offer data sources alongside its set of resource types

```
data "aws_s3_bucket" "existing_bucket" {  
    bucket = "dev-chtz-testbucket-c"  
}
```

```
resource "aws_s3_bucket_public_access_block" "bucket_nonpublic" {  
    bucket = data.aws_s3_bucket.existing_bucket.id  
    block_public_acls = true  
}
```



# Terraform

## Input Variables



- Input variables serve as parameters for a Terraform module, allowing aspects of the module to be customized
- You can set the root module variable using CLI options, environment variables or .tfvars files (which are either loaded automatically or must be provided via CLI options)

```
variable "block_public_acls" {  
    type = bool  
    default = true  
}
```

```
resource "aws_s3_bucket_public_access_block" "samplebucket_nonpublic" {  
    count = var.block_public_acls ? 1 : 0  
    bucket = aws_s3_bucket.samplebucket.id  
    block_public_acls = true  
}
```

# Terraform



## Local Values

- A local value assigns a name to an expression, allowing it to be used multiple times within a module without repeating it

```
locals {  
  common_tags = {  
    TFWorkspace = terraform.workspace  
  }  
}  
  
resource "aws_s3_bucket" "samplebucket" {  
  bucket = "${terraform.workspace}-${var.bucket_suffix}"  
  
  tags = merge(local.common_tags, {  
    Suffix : var.bucket_suffix  
  })  
}
```

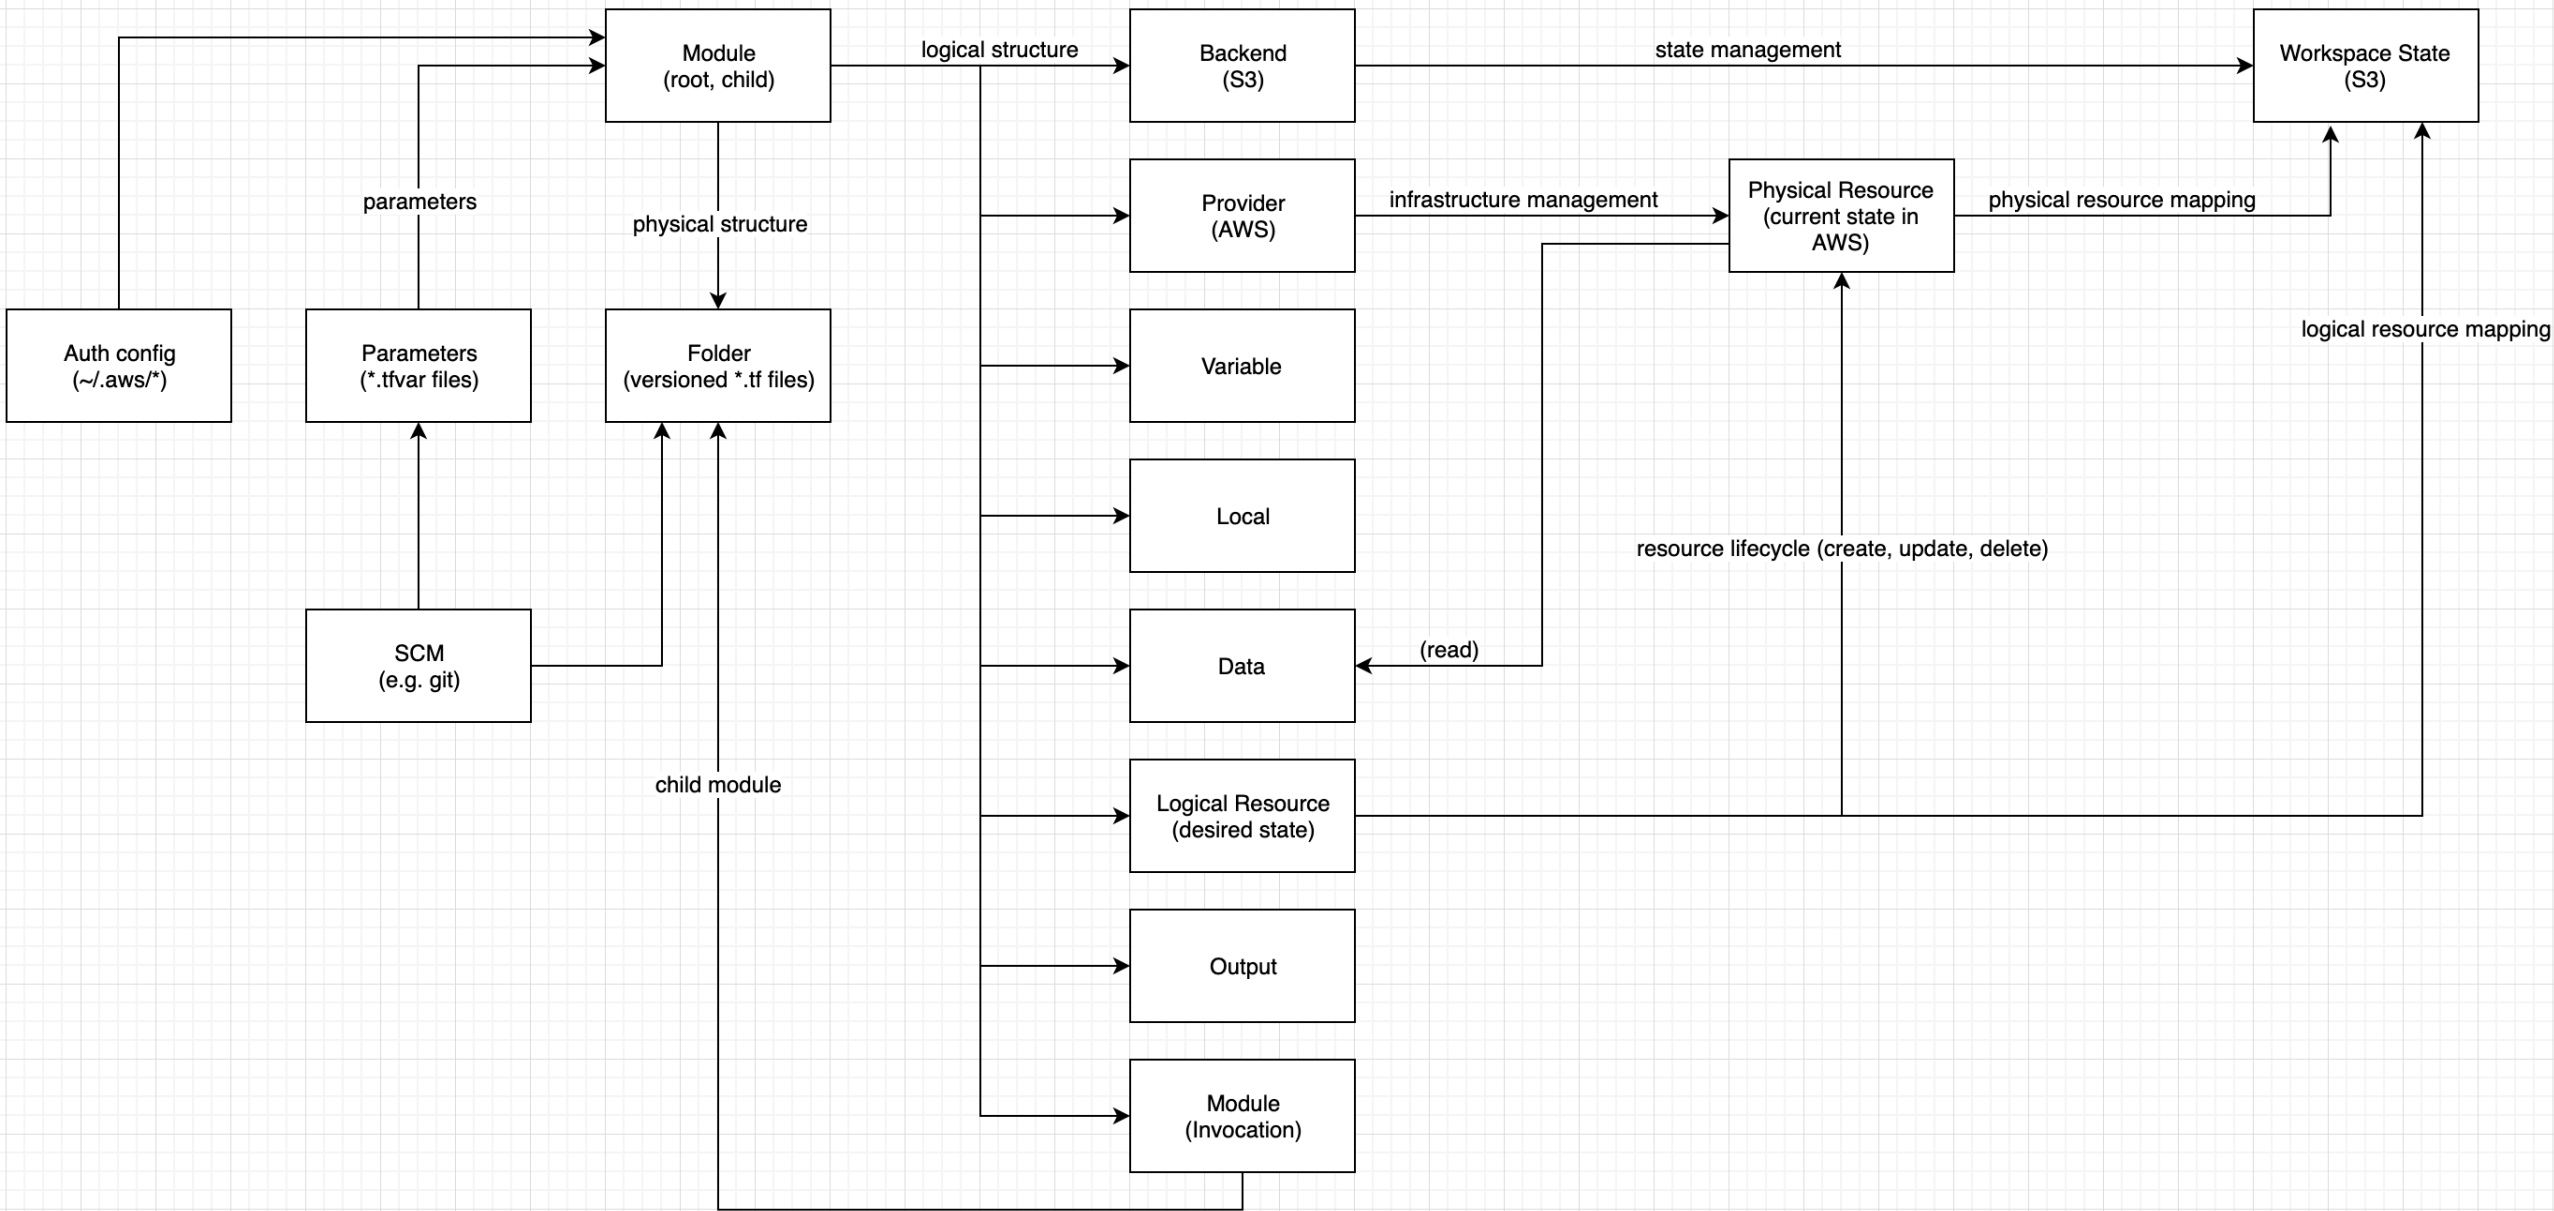
# Terraform

## Output Values



- Output values are like the return values of a Terraform module
- A child module can use outputs to expose a subset of its resource attributes to a parent module
- A root module can use outputs to print certain values in the CLI output after running terraform apply
- When using remote state, root module outputs can be accessed by other configurations via a terraform\_remote\_state data source

```
output "bucket" {  
    value = aws_s3_bucket.samplebucket.id  
}
```



# Templating

# Terraform

## Templating



```
variable "alist" {
  type = "list"
  default = [ {name = "Hallo", value="Welt"}, {name = "hello", value="world"} ]
}

data "template_file" "example" {
  template = <<EOT
{
  "environment":
    [ %{ for i, item in var.alist ~}
      { "name": ${jsonencode(item.name)}, "value": ${jsonencode(item.value)} }
      %{ if i < length(var.alist)-1 ~} , %{endif ~}
    %{ endfor ~} ],
  "better_environment" : ${jsonencode(var.alist)}
}
EOT
}
```

# Workspaces

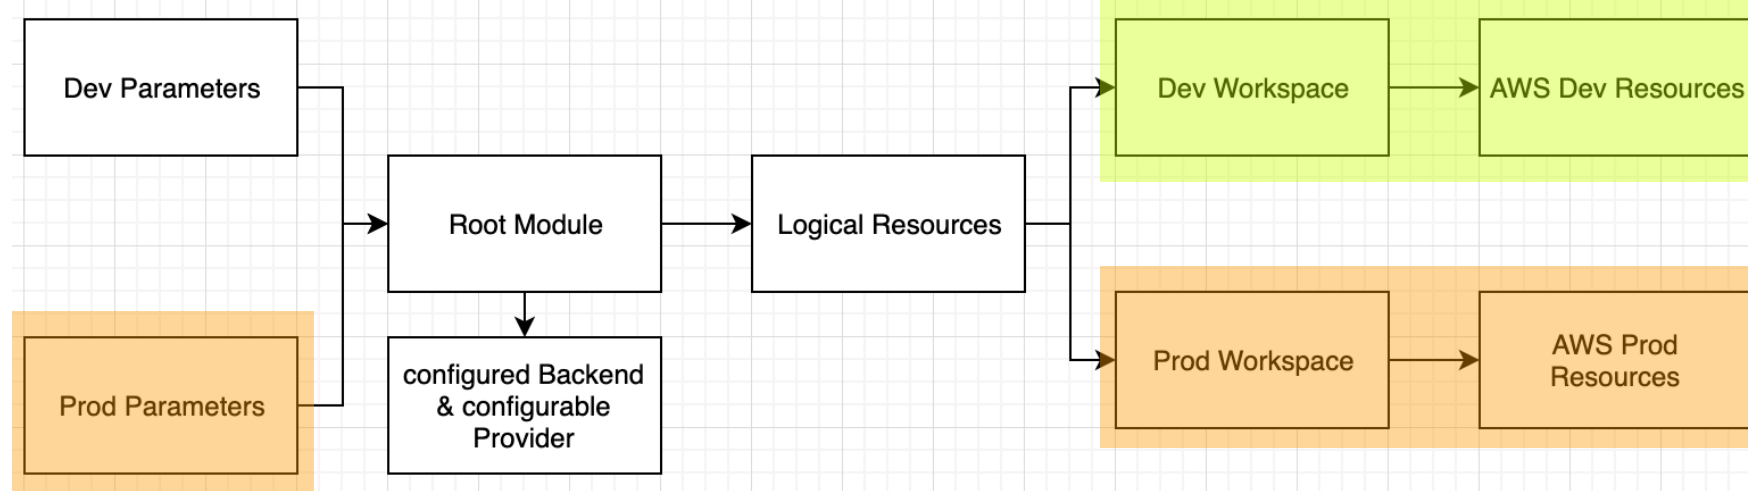
# Terraform

## Workspaces

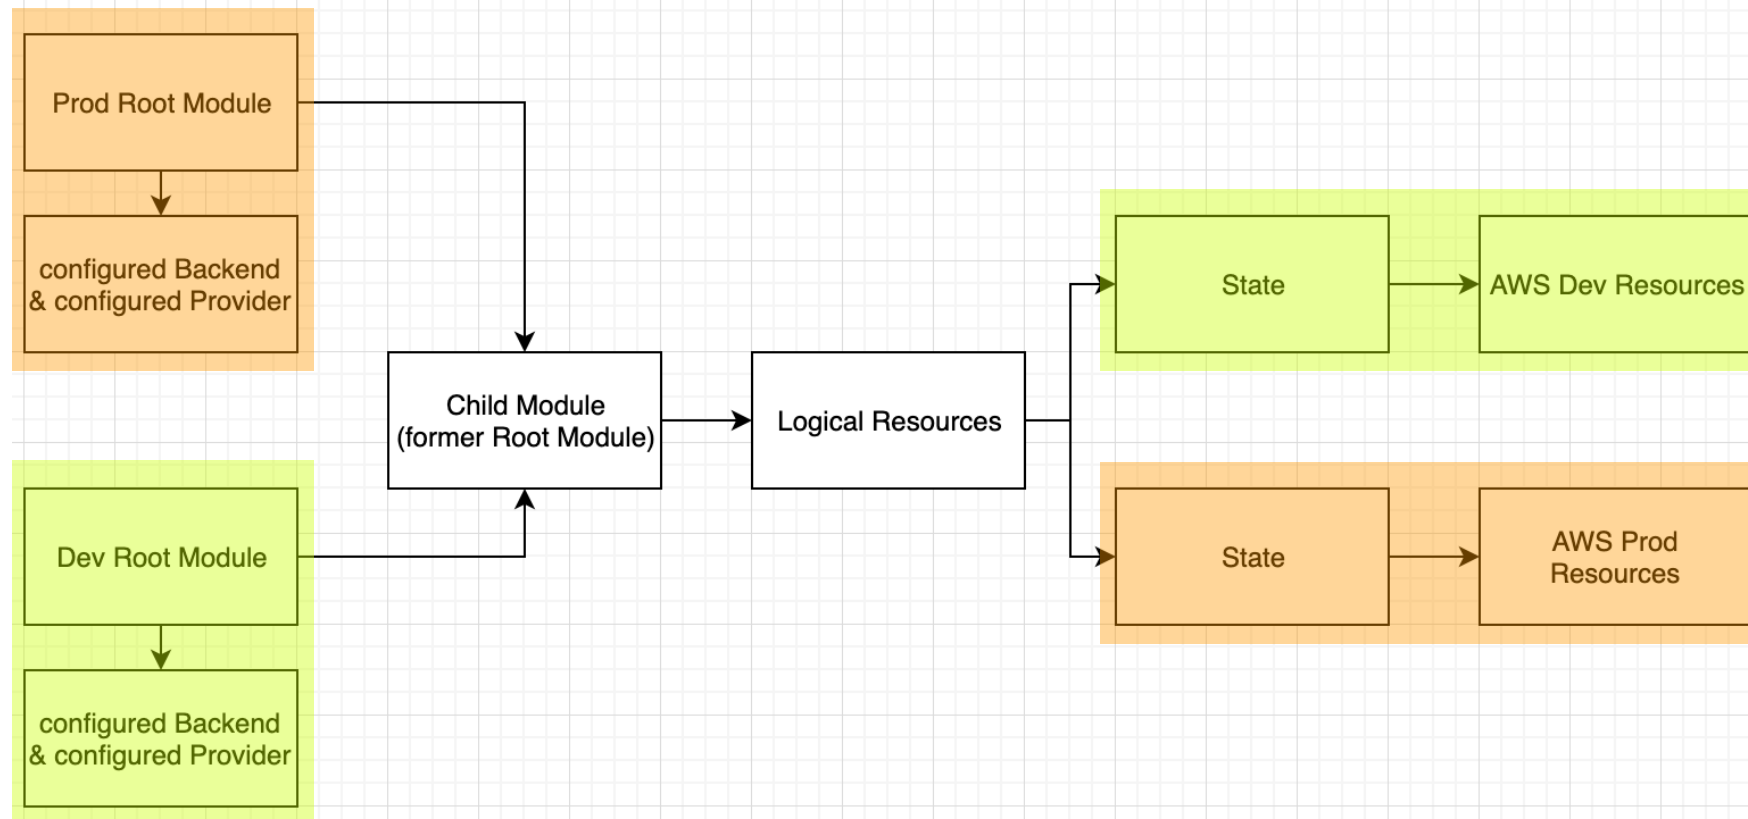


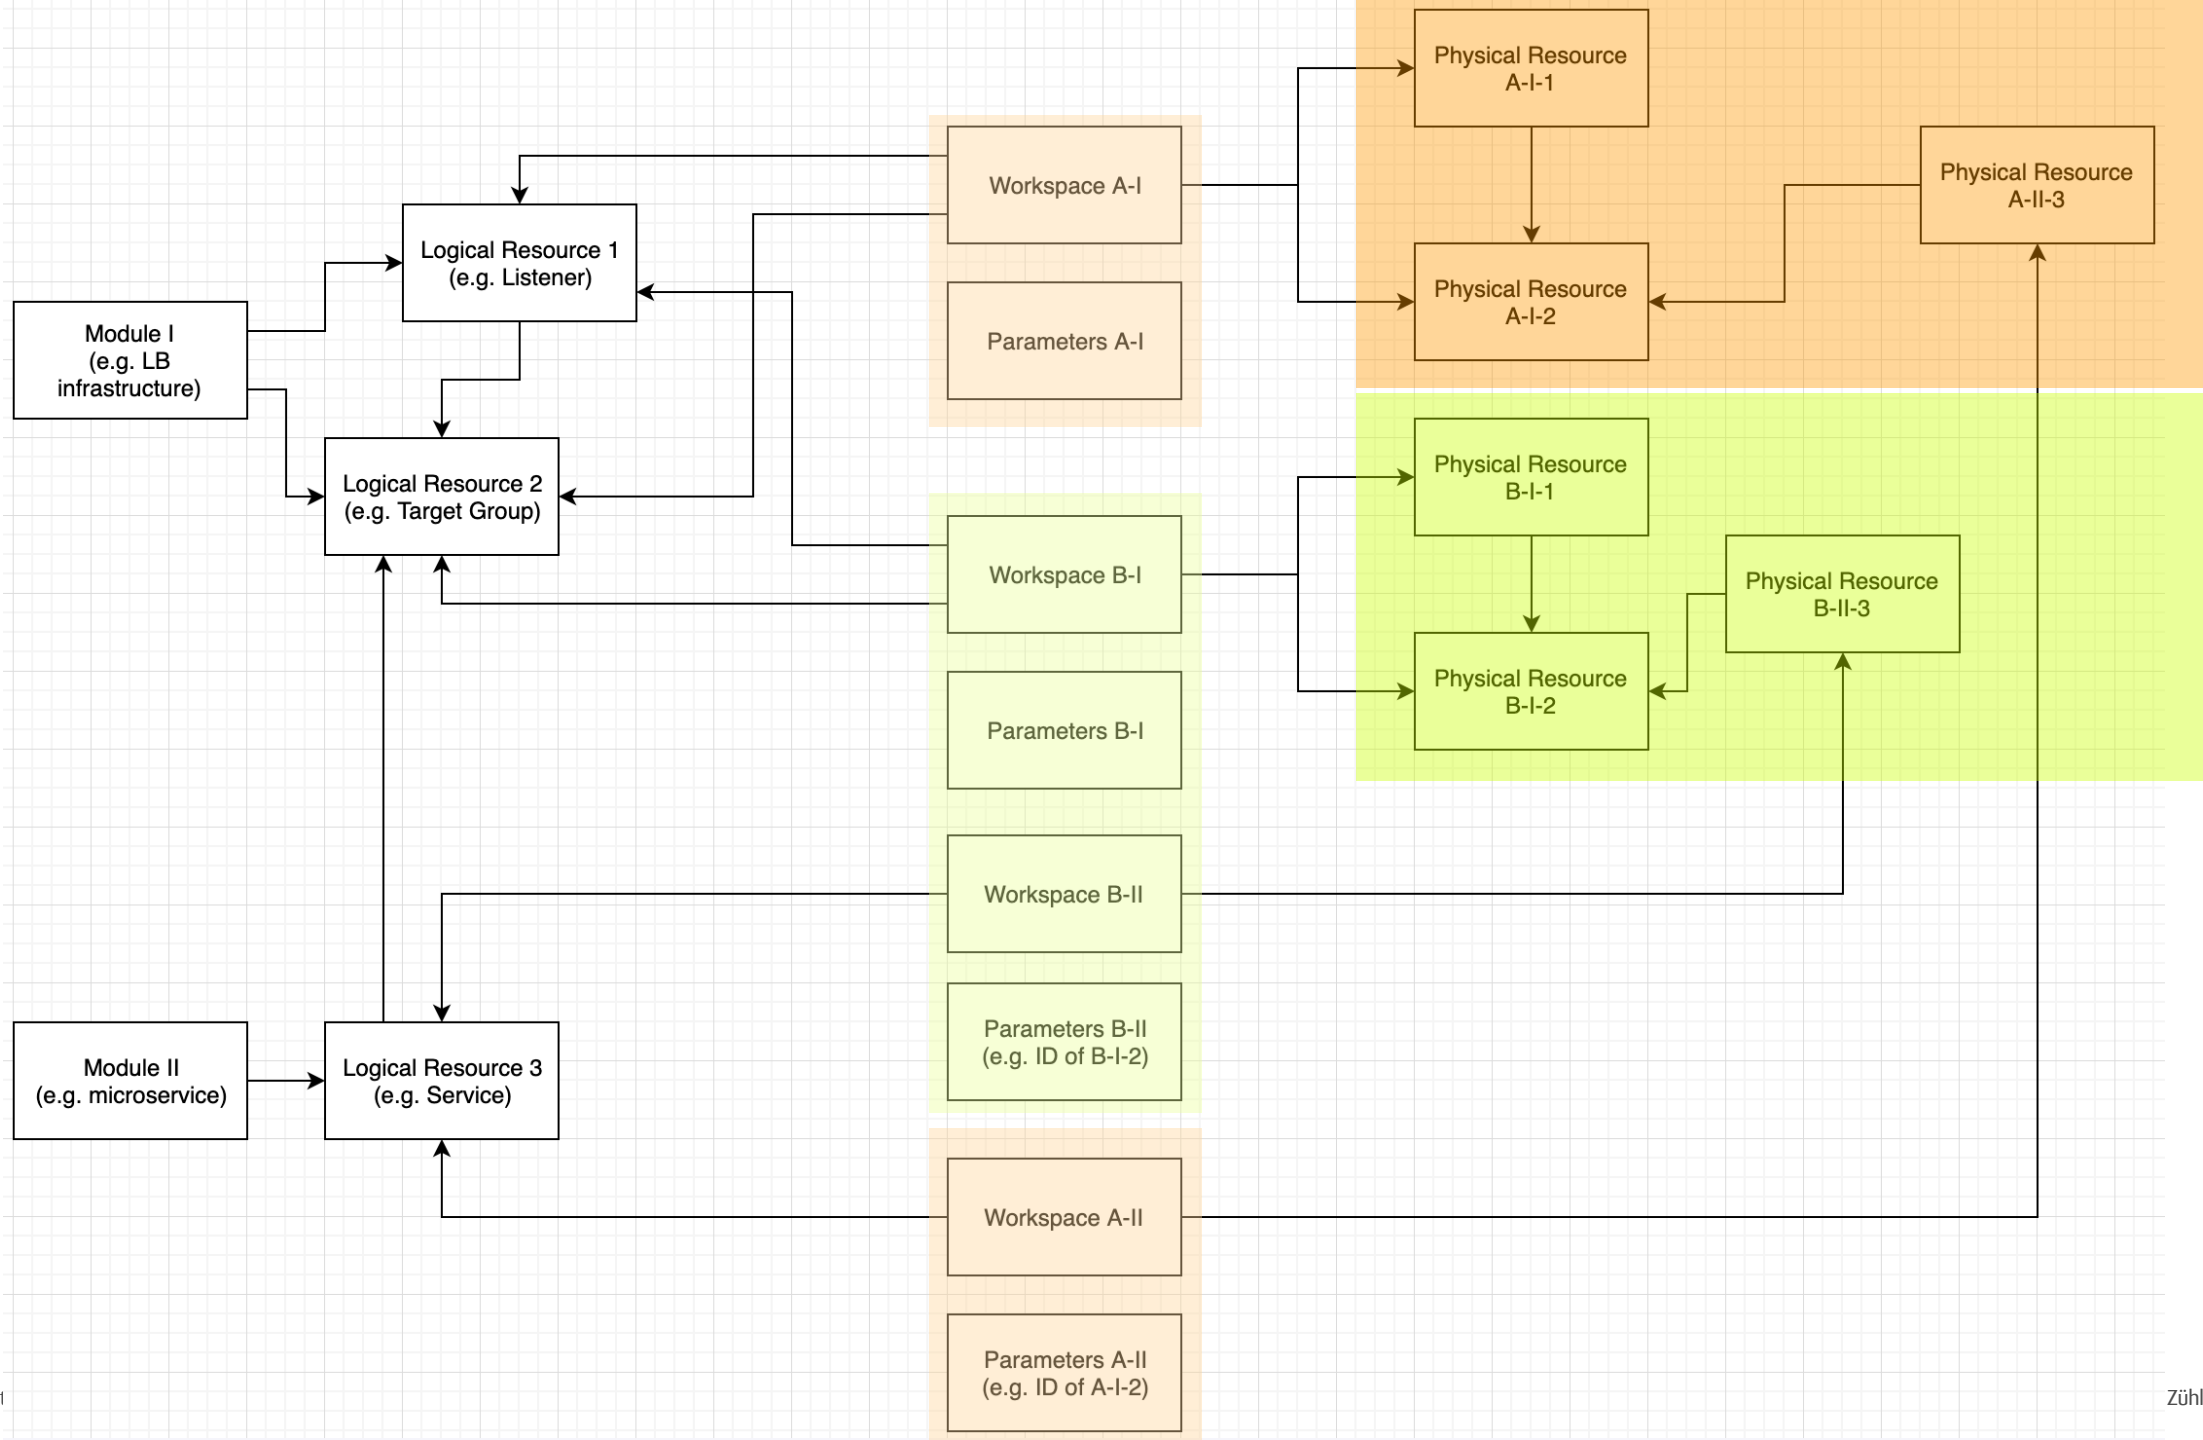
- The S3 backend supports multiple named workspaces, allowing multiple states to be associated with a single configuration
- Within your Terraform configuration, you may include the name of the current workspace using the `${terraform.workspace}` interpolation sequence
- When Terraform is used to manage larger systems, teams should use multiple separate Terraform configurations that correspond with suitable architectural boundaries within the system. Workspaces alone are not a suitable tool for system decomposition.
- The input variable values passed to root modules are not stored in the workspace state.
- Name your workspaces with both their component and their environment (e.g. vpc-dev)





XOR





# Samples

# Samples

see Github



- `bucket sample` – workspaces, variables, backend, provider, resources, tagging
- `childmodules` – modules, breaking and fixing terraform state
- `datasources` – `datasources`
- `multiaccount` – multiple root modules instead of workspaces
- `strings` – JSON and strings

# Q & A